

AI Chat Bot Operations Runbook

Storage source of truth

When `FIRESTORE_SOURCE_OF_TRUTH=true`, startup hydration loads bot configuration, knowledge, and business data into the runtime cache. Orders, appointments, customer profiles, transcripts, escalations, takeover state, and human reply ledgers read from that cache. `FIRESTORE_MIRROR_WRITES=true` updates Firestore after successful local writes and refreshes the cache.

Verification after deployment

Run the full unit suite and deployment regression suite. Verify HTTP 200 health, PM2 online status, and no Firestore hydration or mirror errors. Keep the rollback archive and SHA-256 beside the evidence record.

Backup and restore

Back up data/ outside the VPS without `.env` or provider credentials. Record SHA-256, restore into a temporary directory, count files and bytes, then remove the temporary extraction. Firestore rollback rehearsal is separate.

Logs and human operations

Audit and escalation APIs support limit and cursor. Old JSONL records move to `.archive.jsonl` before the active file is capped. Human replies use idempotency operation IDs; pending or unknown delivery must be reconciled before retrying.

Security boundaries

Client `auth.json` files remain local and are excluded from Git. Keep provider tokens and customer exports out of the repository and deployment payload. Dashboard and human reply routes require account ownership checks.